

assignment5_imagenette

May 5, 2025

1 MP5: Training Your Diffusion Model!

1.1 Setup environment

```
[1]: # Import essential modules. Feel free to add whatever you need.
import matplotlib.pyplot as plt
import torch
from torch import optim
from torch import nn
import torch.nn.functional as F
from torch.utils.data import DataLoader, Subset
from torchvision.datasets import MNIST
from torchvision.transforms import ToTensor
from tqdm.notebook import tqdm # For progress bars
```

1.2 Visualization helper function

```
[2]: def visualize_images_with_titles(images: torch.Tensor, column_names: list[str]):
    """
    Visualize images as a grid and title the columns with the provided names.

    Args:
        images: (N, C, H, W) tensor of images, where N is (number of rows *
        ↪ number of columns)
        column_names: List of column names for the titles.

    Example usage:
        visualize_images_with_titles(torch.randn(16, 1, 32, 32), ['1', '2', '3',
        ↪ '4'])
    """
    num_images, num_columns = images.shape[0], len(column_names)
    assert num_images % num_columns == 0, 'Number of images must be a multiple
    ↪ of the number of columns.'

    num_rows = num_images // num_columns
    fig, axes = plt.subplots(num_rows, num_columns, figsize=(num_columns * 1,
    ↪ num_rows * 1))
```

```

for i, ax in enumerate(axes.flat):
    img = images[i].permute(1, 2, 0).cpu().detach().numpy()
    ax.imshow(img, cmap='gray')
    ax.axis('off')
    if i < num_columns:
        ax.set_title(column_names[i % num_columns])

plt.tight_layout()
plt.show()

```

2 Part 1: Training a Single-step Denoising UNet

2.1 Implementing Simple and Composed Ops

```

[3]: class Conv(nn.Module):
    """
    Conv is a convolutional layer that doesn't change the image resolution,
    only the channel dimension.
    """
    def __init__(self, in_channels: int, out_channels: int):
        super().__init__()
        self.in_channels = in_channels
        self.out_channels = out_channels

        self.conv = nn.Conv2d(in_channels=in_channels,
    out_channels=out_channels, kernel_size=1)
        self.bn = nn.BatchNorm2d(num_features=out_channels)
        self.gelu = nn.GELU()

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        """
        Conv2d(3,1,1) + BN (nn.BatchNorm2d()) + GELU (nn.GELU())
        """
        conv_result = self.conv(x)
        bn_result = self.bn(conv_result)
        gelu_result = self.gelu(bn_result)
        assert gelu_result.shape == (x.shape[0], self.out_channels, x.shape[2],
    x.shape[3]), \
            f"Expected shape {(x.shape[0], self.out_channels, x.shape[2], x.
    shape[3])}, but got {gelu_result.shape}"
        return gelu_result

class DownConv(nn.Module):
    """

```

```

DownConv is a convolutional layer that downsamples the tensor by 2.
D_in, H, W => D_out, H/2, W/2
"""
def __init__(self, in_channels: int, out_channels: int):
    super().__init__()
    self.in_channels = in_channels
    self.out_channels = out_channels

    self.conv = nn.Conv2d(in_channels=in_channels,
↪out_channels=out_channels, kernel_size=3, stride=2, padding=1)
    self.bn = nn.BatchNorm2d(num_features=out_channels)
    self.gelu = nn.GELU()

def forward(self, x: torch.Tensor) -> torch.Tensor:
    """Conv2d (3,2,1) + BN + GELU"""
    # Conv2d(3,2,1) + BN (nn.BatchNorm2d()) + GELU (nn.GELU())
    conv_result = self.conv(x)
    bn_result = self.bn(conv_result)
    gelu_result = self.gelu(bn_result)
    assert gelu_result.shape == (x.shape[0], self.out_channels, x.shape[2] /
↪2, x.shape[3] // 2), \
        f"Expected shape {(x.shape[0], self.out_channels, x.shape[2] // 2,
↪x.shape[3] // 2)}, but got {gelu_result.shape}"
    return gelu_result

class UpConv(nn.Module):
    """
    UpConv is a convolutional layer that upsamples the tensor by 2.
    """
    def __init__(self, in_channels: int, out_channels: int):
        super().__init__()
        self.in_channels = in_channels
        self.out_channels = out_channels
        self.conv_transpose = nn.ConvTranspose2d(in_channels=in_channels,
↪out_channels=out_channels, kernel_size=4, stride=2, padding=1)
        self.bn = nn.BatchNorm2d(num_features=out_channels)
        self.gelu = nn.GELU()

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        """ConTranspose2d (4,2,1) + BN + GELU"""
        # ConvTranspose2d(4,2,1) + BN (nn.BatchNorm2d()) + GELU (nn.GELU())
        conv_result = self.conv_transpose(x)
        bn_result = self.bn(conv_result)
        gelu_result = self.gelu(bn_result)
        assert gelu_result.shape == (x.shape[0], self.out_channels, x.shape[2]
↪* 2, x.shape[3] * 2), \

```

```

        f"Expected shape {(x.shape[0], self.out_channels, x.shape[2] * 2, x.
↳shape[3] * 2)}, but got {gelu_result.shape}"
        return gelu_result

class Flatten(nn.Module):
    """
    Flatten is an average pooling layer that flattens a 7x7 tensor into a 1x1
↳tensor.
    7 is the resulting height and width after the downsampling operations.
    """

    def __init__(self):
        super().__init__()
        self.avg_pool = nn.AvgPool2d(kernel_size=7)
        self.gelu = nn.GELU()

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        """AvgPool (7) + GELU"""
        # AvgPool(7) + GELU (nn.GELU())
        avg_pool_result = self.avg_pool(x)
        gelu_result = self.gelu(avg_pool_result)
        assert gelu_result.shape == (x.shape[0], x.shape[1], 1, 1), 'GELU
↳output shape is incorrect'
        return gelu_result

class Unflatten(nn.Module):
    """
    Unflatten is a convolutional layer that unflattens/upsamples a 1x1 tensor
↳into a 7x7 tensor.
    """

    def __init__(self, in_channels: int):
        super().__init__()
        self.in_channels = in_channels
        self.conv_transpose = nn.ConvTranspose2d(in_channels=in_channels,
↳out_channels=in_channels, kernel_size=7, stride=1, padding=0)
        self.bn = nn.BatchNorm2d(num_features=in_channels)
        self.gelu = nn.GELU()

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        """ConvTranspose2d(7,7,0) + BN + GELU"""
        # ConvTranspose2d(7,7,0) + BN (nn.BatchNorm2d()) + GELU (nn.GELU())
        conv_result = self.conv_transpose(x)
        bn_result = self.bn(conv_result)
        gelu_result =self.gelu(bn_result)

```

```

        assert gelu_result.shape == (x.shape[0], self.in_channels, 7, 7), \
            f"Expected shape {(x.shape[0], self.in_channels, 7, 7)}, but got {
↳gelu_result.shape}"
        return gelu_result

class ConvBlock(nn.Module):
    """
    similar to Conv but includes an additional Conv.
    Note that it has the same input and output shape as (1) Conv.
    """
    def __init__(self, in_channels: int, out_channels: int):
        super().__init__()
        self.in_channels = in_channels
        self.out_channels = out_channels
        self.conv1 = nn.Conv2d(in_channels=in_channels,
↳out_channels=out_channels, kernel_size=1)
        self.bn1 = nn.BatchNorm2d(num_features=out_channels)
        self.gelu1 = nn.GELU()
        self.conv2 = nn.Conv2d(in_channels=out_channels,
↳out_channels=out_channels, kernel_size=1)
        self.bn2 = nn.BatchNorm2d(num_features=out_channels)
        self.gelu2 = nn.GELU()

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        """
        Conv2d(3,1,1) + BN (nn.BatchNorm2d()) + GELU (nn.GELU()) +
↳Conv2d(1,1,1) + BN (nn.BatchNorm2d()) + GELU (nn.GELU())
        """
        conv_result = self.conv1(x)
        bn_result = self.bn1(conv_result)
        gelu_result = self.gelu1(bn_result)
        conv_result_2 = self.conv2(gelu_result)
        bn_result_2 = self.bn2(conv_result_2)
        gelu_result_2 = self.gelu2(bn_result_2)

        # Check if the output shape is as expected
        assert gelu_result_2.shape == (x.shape[0], self.out_channels, x.
↳shape[2], x.shape[3]), \
            f"Expected shape {(x.shape[0], self.out_channels, x.shape[2], x.
↳shape[3])}, but got {gelu_result_2.shape}"
        return gelu_result_2

class DownBlock(nn.Module):
    """

```

*similar to DownConv but includes an additional ConvBlock.
Note that it has the same input and output shape as (2) DownConv.*

```

"""
def __init__(self, in_channels: int, out_channels: int):
    super().__init__()
    self.in_channels = in_channels
    self.out_channels = out_channels

    self.down_conv = DownConv(in_channels, out_channels)
    self.conv_block = ConvBlock(out_channels, out_channels)

def forward(self, x: torch.Tensor) -> torch.Tensor:
    """DownConv + ConvBlock"""
    # DownConv + ConvBlock
    downconv_result = self.down_conv(x)
    convblock_result = self.conv_block(downconv_result)

    # Check if the output shape is as expected
    assert convblock_result.shape == (x.shape[0], self.out_channels, x.
↪shape[2] // 2, x.shape[3] // 2), \
        f"Expected shape {(x.shape[0], self.out_channels, x.shape[2] // 2,
↪x.shape[3] // 2)}, but got {convblock_result.shape}"

    return convblock_result

```

```

class UpBlock(nn.Module):
    """
    similar to UpConv but includes an additional ConvBlock.
    Note that it has the same input and output shape as (3) UpConv.
    """

    def __init__(self, in_channels: int, out_channels: int):
        super().__init__()
        self.in_channels = in_channels
        self.out_channels = out_channels

        self.up_conv = UpConv(in_channels, out_channels)
        self.conv_block = ConvBlock(out_channels, out_channels)

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        """UpConv + ConvBlock"""
        # UpConv + ConvBlock
        upconv_result = self.up_conv(x)
        convblock_result = self.conv_block(upconv_result)

        # Check if the output shape is as expected

```

```

        assert convblock_result.shape == (x.shape[0], self.out_channels, x.
↪shape[2] * 2, x.shape[3] * 2), \
            f"Expected shape {(x.shape[0], self.out_channels, x.shape[2] * 2, x.
↪shape[3] * 2)}, but got {convblock_result.shape}"
        return convblock_result

```

2.2 Implementing Unconditional UNet

```

[4]: class UnconditionalUNet(nn.Module):
    def __init__(
        self,
        in_channels: int,
        num_hiddens: int,
    ):
        self.in_channels = in_channels
        self.num_hiddens = num_hiddens

        super().__init__()

        # Layers in U-Net
        self.first_conv = ConvBlock(self.in_channels, num_hiddens)

        # Downsampling Path
        self.down1 = DownBlock(num_hiddens, num_hiddens)
        self.down2 = DownBlock(num_hiddens, num_hiddens * 2)

        # Flatten and Unflatten (bottom of U-Net)
        self.flatten = Flatten()
        self.unflatten = Unflatten(num_hiddens * 2)

        # Upsampling Path
        # upsample1 is the one that receives the concatenatio of Unflatten and
↪DownBlock
        self.upsample1 = UpBlock(num_hiddens * 4, num_hiddens)
        self.upsample2 = UpBlock(num_hiddens * 2, num_hiddens)

        # Final Convolution Blocks
        self.final_conv1 = ConvBlock(num_hiddens * 2, num_hiddens)
        self.final_conv2 = nn.Conv2d(num_hiddens, self.in_channels,
↪kernel_size=3, stride=1, padding=1)

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        assert x.shape[-2:] == (28, 28), "Expect input shape to be (28, 28)."

        # Encoder?
        c1 = self.first_conv(x)
        d1 = self.down1(c1)

```

```

        d2 = self.down2(d1)

        # Bottleneck
        flat = self.flatten(d2)
        b = self.unflatten(flat)

        # Decoder
        cat1 = torch.cat((b, d2), dim=1)
        u1 = self.upsample1(cat1)

        ## Second level
        cat2 = torch.cat((u1, d1), dim=1)
        u2 = self.upsample2(cat2)

        ## Final (top) level
        cat3 = torch.cat((u2, c1), dim=1)
        f1 = self.final_conv1(cat3)

        # Final convolution
        out_res = self.final_conv2(f1)

        assert out_res.shape == x.shape, f"Output shape {out_res.shape} doesn't match input shape {x.shape}"

        return out_res

```

2.3 Visualizing the noising process

```

[5]: dataset = MNIST(root="data", download=True, transform=ToTensor(), train=True)


train_loader = DataLoader(dataset, batch_size=16, shuffle=True)
# Get a batch of data
for images, labels in train_loader:
    break
# Create an instance of the model
model = UnconditionalUNet(in_channels=1, num_hiddens=16)


# Forward pass
output = model(images)
# Visualize the output
visualize_images_with_titles(output, ['Output 1', 'Output 2', 'Output 3',
    'Output 4'])
# Visualize the input
visualize_images_with_titles(images, ['Input 1', 'Input 2', 'Input 3', 'Input
    4'])

```



```
# Visualize the labels by printing them
print("Labels (4x4 matrix):")
print(labels.reshape(4, 4))
```





Labels (4x4 matrix):
 tensor([[1, 5, 7, 7],
 [2, 0, 4, 6],
 [8, 2, 7, 9],
 [0, 0, 2, 6]])

```
[6]: ITERATIONS = 9
      NUM_IMAGES = 5

      dataset = MNIST(root="data", download=True, transform=ToTensor(), train=True)
      train_loader = DataLoader(dataset, batch_size=16, shuffle=True)

      for images, labels in train_loader:
          break
      # Create an instance of the model
      model = UnconditionalUNet(in_channels=1, num_hiddens=16)

      with torch.no_grad():
          # Select the first {NUM_IMAGES} images in the batch
          imgs = images[:NUM_IMAGES]
          progression = [imgs.clone()]

          # Pass the images through the model for {ITERATIONS} iterations
```


```

for i in range(ITERATIONS):
    imgs = model(imgs)
    progression.append(imgs.clone())

# Plot the original images and their progression
num_iters = len(progression) # 9 iterations (initial + ITERATIONS steps)
num_imgs = progression[0].shape[0] # NUM_IMAGES images

fig, axes = plt.subplots(num_imgs, num_iters, figsize=(num_iters * 2, num_imgs * 2))
for i in range(num_imgs):
    for j in range(num_iters):
        axes[i, j].imshow(progression[j][i].squeeze().cpu(), cmap="gray")
        axes[i, j].axis("off")
        if i == 0:
            axes[i, j].set_title(f"Iter {j}")
plt.tight_layout()
plt.show()

```



2.4 Training a Single-Step Unconditional UNet

- Plot the loss curve
- Sample results on the test set

```

[7]: dataset = MNIST(root='data', download=True, transform=ToTensor(), train=True)
    dataloader = DataLoader(dataset, batch_size=256, shuffle=True)

```

```

EPOCHS = 5
BATCH_SIZE = 256
LEARNING_RATE = 1e-4
NUM_HIDDENS = 128 # 'D' in the UNet diagram
NOISE_SIGMA = 0.5 # Fixed noise level for this part
DEVICE = torch.device("cuda" if torch.cuda.is_available() else "mps" if torch.
    ↪mps.is_available() else "cpu")
print(f"Using device: {DEVICE}")

```

Using device: mps

```

[8]: # DataLoaders
train_dataset = MNIST(root='data', download=True, transform=ToTensor(), ↵
    ↪train=True)
test_dataset = MNIST(root='data', download=True, transform=ToTensor(), ↵
    ↪train=False)

train_dataloader = DataLoader(train_dataset, batch_size=BATCH_SIZE, ↵
    ↪shuffle=True)
# We only need a small fixed batch from the test set for visualization

test_subset_indices = range(8) # Let's visualize 20 images
test_subset = Subset(test_dataset, test_subset_indices)
vis_dataloader = DataLoader(test_subset, batch_size=len(test_subset_indices)) # ↵
    ↪Load all images at once

# Initialization of the model + move to device
model = UnconditionalUNet(in_channels=1, num_hiddens=NUM_HIDDENS).to(DEVICE)

criterion = nn.MSELoss() # L2 Loss
optimizer = optim.Adam(model.parameters(), lr=LEARNING_RATE)

# --- 3. Visualization Prep ---\
# Get the fixed batch of test images
fixed_vis_images, _ = next(iter(vis_dataloader))
fixed_vis_images = fixed_vis_images.to(DEVICE)
# Generate fixed noise for these images ONCE
torch.manual_seed(42) # for reproducible noise
fixed_noise = torch.randn_like(fixed_vis_images).to(device=DEVICE) * NOISE_SIGMA

# Create the fixed noisy versions (clamp for visualization only)
fixed_noisy_vis_images = torch.clamp(fixed_vis_images + fixed_noise, 0, 1)

[9]: # --- 4. Training Loop ---\
train_losses = []
step_count = 0

```

```

for epoch in range(EPOCHS):
    model.train() # Set model to training mode
    epoch_loss = 0.0
    progress_bar = tqdm(train_dataloader, desc=f"Epoch {epoch+1}/{EPOCHS}",
        ↪leave=False)

    for batch_idx, (clean_images, _) in enumerate(progress_bar):
        clean_images = clean_images.to(DEVICE) # (N, 1, 28, 28)

        # --- Generate Noisy Input for this batch ---\
        # Generate NEW noise for every batch
        noise = torch.randn_like(clean_images) * NOISE_SIGMA
        noisy_input = clean_images + noise
        # Note: We don't clamp the input to the model, only the visualized noise
        noisy_input = noisy_input.to(DEVICE) # (N, 1, 28, 28)

        denoised_output = model(noisy_input)

        # Loss is between model output and the ORIGINAL CLEAN image
        loss = criterion(denoised_output, clean_images)

        # --- Backward Pass & Optimization ---\
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # --- Logging ---\
        current_loss = loss.item()
        epoch_loss += current_loss
        train_losses.append(current_loss)
        step_count += 1
        progress_bar.set_postfix(loss=f"{current_loss:.4f}")

    avg_epoch_loss = epoch_loss / len(train_dataloader)
    print(f"Epoch {epoch+1}/{EPOCHS} - Average Training Loss: {avg_epoch_loss:.
        ↪4f}")

    # Visualization for epoch 1 and 5
    if (epoch + 1) == 1 or (epoch + 1) == EPOCHS:
        print(f"\n--- Visualizing after Epoch {epoch+1} ---")
        model.eval()
        with torch.no_grad():
            denoised_test_output = model(fixed_vis_images + fixed_noise)

        vis_clean = fixed_vis_images.cpu()

```

```

vis_noisy = fixed_noisy_vis_images.cpu()
vis_denoised = torch.clamp(denoised_test_output.cpu(), 0, 1)

num_vis = len(vis_clean)
images_to_show = torch.stack([\
    vis_clean.squeeze(1),\
    vis_noisy.squeeze(1),\
    vis_denoised.squeeze(1)\
]).permute(1, 0, 2, 3).reshape(num_vis * 3, 1, 28, 28) # Reshape for 
↪grid

visualize_images_with_titles(
    images_to_show,
    column_names=['Input', f'Noisy (={NOISE_SIGMA})', 'Denoised']
)
print("-" * 20)

plt.figure(figsize=(10, 5))
plt.plot(range(step_count), train_losses)
plt.xlabel("Training Step")
plt.ylabel("MSE Loss")
plt.title("Training Loss Curve (Single-Step Denoising)")
plt.yscale('log') # Often helpful to see loss changes
plt.grid(True)
plt.show()


print("Training finished.")

```

Epoch 1/5: 0%| | 0/235 [00:00<?, ?it/s]

Epoch 1/5 - Average Training Loss: 0.0247

--- Visualizing after Epoch 1 ---




Epoch 2/5: 0%|


| 0/235 [00:00<?, ?it/s]

```
Epoch 2/5 - Average Training Loss: 0.0114
Epoch 3/5:  0%|          | 0/235 [00:00<?, ?it/s]
Epoch 3/5 - Average Training Loss: 0.0100
Epoch 4/5:  0%|          | 0/235 [00:00<?, ?it/s]
Epoch 4/5 - Average Training Loss: 0.0094
Epoch 5/5:  0%|          | 0/235 [00:00<?, ?it/s]
Epoch 5/5 - Average Training Loss: 0.0091

--- Visualizing after Epoch 5 ---
```


Input Noisy ($\sigma=0.5$) Denoised





Training finished.

```
[10]: PATH = "unconditional_unet.pth"
torch.save(model.state_dict(), PATH)
print(f"Model saved as {PATH}")
```

Model saved as unconditional_unet.pth

2.5 Out-of-Distribution Testing

```
[11]: PATH = "unconditional_unet.pth"
model = UnconditionalUNet(in_channels=1, num_hiddens=NUM_HIDDENS).to(DEVICE)
model.load_state_dict(torch.load(PATH, map_location=DEVICE))
model.eval()
print("Model loaded from", PATH)
```

Model loaded from unconditional_unet.pth

```
[12]: test_images = next(iter(vis_dataloader))[0].to(DEVICE)
sigmas = [0.0, 0.2, 0.4, 0.5, 0.6, 0.8, 1.0]
# Select the first image from the fixed batch of test images
img = test_images[0:1]

noisy_images_list = []
denoised_images_list = []

model.eval()
with torch.no_grad():
```

```

for sigma in sigmas:
    noise = sigma * torch.randn_like(img)
    noisy = img + noise
    denoised = model(noisy)
    noisy_images_list.append(noisy.cpu())
    denoised_images_list.append(denoised.cpu())


# Create a grid with 2 rows: first row for noisy, second row for denoised images
fig, axes = plt.subplots(2, len(sigmas), figsize=(len(sigmas) * 2, 4))

for i, sigma in enumerate(sigmas):
    # Top row: noisy image
    ax = axes[0, i]
    ax.imshow(noisy_images_list[i].squeeze(), cmap="gray")
    ax.axis("off")
    ax.set_title(f"={sigma}", fontsize=10)

    # Bottom row: denoised image
    ax = axes[1, i]
    ax.imshow(denoised_images_list[i].squeeze(), cmap="gray")
    ax.axis("off")
    if i == 0:
        ax.set_ylabel("Denoised", fontsize=10)

plt.tight_layout()
plt.show()

```



3 Part 2: Training a Diffusion Model

3.1 Implementing a Time-conditioned UNet

```

[13]: class FCBlock(nn.Module):
    def __init__(self, in_channels: int, out_channels: int):
        super().__init__()

```

```

        self.in_channels = in_channels
        self.out_channels = out_channels

        # Definitions for layer
        self.fc = nn.Linear(in_channels, out_channels)
        self.gelu = nn.GELU()

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        x = self.fc(x)
        x = self.gelu(x)
        return x


class TimeConditionalUNet(nn.Module):
    def __init__(
        self,
        in_channels: int,
        num_classes: int,
        num_hiddens: int,
    ):
        super().__init__()
        self.in_channels = in_channels
        self.num_classes = num_classes
        self.num_hiddens = num_hiddens

        # Time embedding
        time_embed_dim = num_hiddens * 4
        self.time_embed = nn.Sequential(
            FCBlock(1, time_embed_dim),
            FCBlock(time_embed_dim, time_embed_dim)
        )

        # Projection layers for time embedding
        self.time_proj_1 = nn.Linear(time_embed_dim, num_hiddens)
        self.time_proj_2 = nn.Linear(time_embed_dim, num_hiddens)
        self.time_proj_3 = nn.Linear(time_embed_dim, num_hiddens * 2)

        # Encoder path
        self.first_conv = ConvBlock(in_channels, num_hiddens)
        self.down1 = DownBlock(num_hiddens, num_hiddens)
        self.down2 = DownBlock(num_hiddens, num_hiddens * 2)

        # Bottleneck
        self.flatten = Flatten()
        self.unflatten = Unflatten(num_hiddens * 2)

```

```

    # Decoder path
    self.upsample1 = UpBlock(num_hiddens * 4, num_hiddens)
    self.upsample2 = UpBlock(num_hiddens * 2, num_hiddens)

    # Final layers
    self.final_conv1 = ConvBlock(num_hiddens * 2, num_hiddens)
    self.final_conv2 = nn.Conv2d(num_hiddens, in_channels, kernel_size=3,
    ↪padding=1)

    def forward(self, x: torch.Tensor, t: torch.Tensor) -> torch.Tensor:
        """
        Args:
            x: (N, C, H, W) input tensor
            t: (N,) normalized time tensor
        """
        assert x.shape[-2:] == (28, 28), "Expected input shape to be (28, 28)"

        # Time embedding
        t = t.unsqueeze(-1).float() # (N, 1)
        t_emb = self.time_embed(t) # (N, time_embed_dim)

        # Project time embeddings for different layers
        t1 = self.time_proj_1(t_emb).unsqueeze(-1).unsqueeze(-1) # (N,
    ↪num_hiddens, 1, 1)
        t2 = self.time_proj_2(t_emb).unsqueeze(-1).unsqueeze(-1) # (N,
    ↪num_hiddens, 1, 1)
        t3 = self.time_proj_3(t_emb).unsqueeze(-1).unsqueeze(-1) # (N,
    ↪num_hiddens * 2, 1, 1)

        # Encoder
        x1 = self.first_conv(x)
        x1 = x1 + t1.expand_as(x1)

        x2 = self.down1(x1)
        x2 = x2 + t2.expand_as(x2)

        x3 = self.down2(x2)
        x3 = x3 + t3.expand_as(x3)

        # Bottleneck
        x_flat = self.flatten(x3)
        x_bottle = self.unflatten(x_flat)

        # Decoder with skip connections
        x = torch.cat([x_bottle, x3], dim=1)
        x = self.upsample1(x)

```

```

x = torch.cat([x, x2], dim=1)
x = self.upsample2(x)

x = torch.cat([x, x1], dim=1)
x = self.final_conv1(x)
x = self.final_conv2(x)

return x

```

3.2 Implementing DDPM Forward and Inverse Process for Time-conditioned Denoising

```

[14]: def ddpm_schedule(beta1: float, beta2: float, num_ts: int) -> dict:
    """Constants for DDPM training and sampling.

    Arguments:
        beta1: float, starting beta value.
        beta2: float, ending beta value.
        num_ts: int, number of timesteps.

    Returns:
        dict with keys:
            betas: linear schedule of betas from beta1 to beta2.
            alphas: 1 - betas.
            alpha_bars: cumulative product of alphas.
    """
    assert beta1 < beta2 < 1.0, "Expect beta1 < beta2 < 1.0."

    # Linear schedule of beta values
    betas = torch.linspace(beta1, beta2, num_ts)

    # Calculate alphas
    alphas = 1 - betas

    # Calculate cumulative product of alphas
    alpha_bars = torch.cumprod(alphas, dim=0)

    return {
        'betas': betas,
        'alphas': alphas,
        'alpha_bars': alpha_bars
    }

```

```

[15]: def ddpm_forward(
    unet: TimeConditionalUNet,
    ddpm_schedule: dict,
    x_0: torch.Tensor,

```

```

    num_ts: int,
) -> torch.Tensor:
    """Algorithm 1 of the DDPM paper.

    Args:
        unet: TimeConditionalUNet
        ddpm_schedule: dict
        x_0: (N, C, H, W) input tensor.
        num_ts: int, number of timesteps.
    Returns:
        (,) diffusion loss.
    """
    unet.train()

    # Sample timestep uniformly
    t = torch.randint(0, num_ts, (x_0.shape[0],), device=x_0.device)

    # Get noise
    epsilon = torch.randn_like(x_0)

    # Get alpha_bar for the sampled timestep
    alpha_bar = ddpm_schedule['alpha_bars'][t].view(-1, 1, 1, 1)

    # Create noisy sample
    x_t = torch.sqrt(alpha_bar) * x_0 + torch.sqrt(1 - alpha_bar) * epsilon

    # Predict noise
    epsilon_theta = unet(x_t, t / num_ts)

    # Calculate loss
    loss = F.mse_loss(epsilon_theta, epsilon)

    return loss

```

```

[16]: @torch.inference_mode()
def ddpm_sample(
    unet: TimeConditionalUNet,
    ddpm_schedule: dict,
    img_wh: tuple[int, int],
    num_ts: int,
    seed: int = 42,
) -> torch.Tensor:
    """Algorithm 2 of the DDPM paper with classifier-free guidance.

    Args:
        unet: TimeConditionalUNet
        ddpm_schedule: dict

```

```

    img_wh: (H, W) output image width and height.
    num_ts: int, number of timesteps.
    seed: int, random seed.

Returns:
    (N, C, H, W) final sample.
"""
UNET.eval()

torch.manual_seed(seed)

# Start from random noise
x_t = torch.randn(1, unet.in_channels, *img_wh, device=next(unet.
↳ parameters()).device)

for t in reversed(range(num_ts)):
    t_tensor = torch.tensor([t], device=x_t.device)
    epsilon_theta = unet(x_t, t_tensor / num_ts)

    alpha = ddpm_schedule['alphas'][t]
    alpha_bar = ddpm_schedule['alpha_bars'][t]
    beta = ddpm_schedule['betas'][t]

    # Add noise if t > 0
    if t > 0:
        noise = torch.randn_like(x_t)
    else:
        noise = 0

    # Denoising step
    x_t = (1 / torch.sqrt(alpha)) * (
        x_t - (beta / torch.sqrt(1 - alpha_bar)) * epsilon_theta
    ) + torch.sqrt(beta) * noise

return x_t

```

```

[17]: class DDPM(nn.Module):
    def __init__(
        self,
        unet: TimeConditionalUNet,
        betas: tuple[float, float] = (1e-4, 0.02),
        num_ts: int = 300,
        p_uncond: float = 0.1,
    ):
        super().__init__()
        self.unet = unet
        self.num_ts = num_ts

```



```

        self.p_uncond = p_uncond

        self.ddpm_schedule = nn.ParameterDict(ddpm_schedule(betas[0], betas[1],
↳ num_ts))

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        """
        Args:
            x: (N, C, H, W) input tensor.

        Returns:
            (,) diffusion loss.
        """
        return ddpm_forward(
            self.unet, self.ddpm_schedule, x, self.num_ts
        )

    @torch.inference_mode()
    def sample(
        self,
        img_wh: tuple[int, int],
        seed: int = 0,
    ):
        return ddpm_sample(
            self.unet, self.ddpm_schedule, img_wh, self.num_ts, seed
        )

```

3.3 Training the Time-conditioned UNet

- Plot the loss curve
- Sample results on the test set

```

[18]: # Set up for training
# dataset = MNIST(root='data', download=True, transform=ToTensor(), train=True)
# dataloader = DataLoader(dataset, batch_size=128, shuffle=True)

# Setup constants and model
EPOCHS = 20
BATCH_SIZE = 32
LEARNING_RATE = 1e-4
NUM_HIDDENS = 32
NUM_TIMESTEPS = 100
BETA1, BETA2 = 1e-4, 0.02
DEVICE = torch.device("cuda" if torch.cuda.is_available() else "mps" if torch.
↳ backends.mps.is_available() else "cpu")

# Initialize model and optimizer

```

```

model = TimeConditionalUNet(
    in_channels=1,
    num_classes=10,
    num_hiddens=NUM_HIDDENS
).to(DEVICE)

optimizer = optim.Adam(model.parameters(), lr=LEARNING_RATE)

# Create DDPM schedule
ddpm_sched = ddpm_schedule(BETA1, BETA2, NUM_TIMESTEPS)
for k, v in ddpm_sched.items():
    ddpm_sched[k] = v.to(DEVICE)

# Setup data
train_dataset = MNIST(root='data', train=True, transform=ToTensor(),
    ↪download=True)
test_dataset = MNIST(root='data', train=False, transform=ToTensor(),
    ↪download=True)

train_loader = DataLoader(train_dataset, batch_size=BATCH_SIZE, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=BATCH_SIZE, shuffle=False)

```

```

[19]: # Training loop
def train_epoch(model, loader, optimizer, ddpm_sched, num_ts):
    model.train()
    total_loss = 0
    progress_bar = tqdm(loader, desc='Training')

    for batch_idx, (images, _) in enumerate(progress_bar):
        images = images.to(DEVICE)

        optimizer.zero_grad()

        # Forward pass
        loss = ddpm_forward(model, ddpm_sched, images, num_ts)

        # Backward pass
        loss.backward()
        optimizer.step()

        # Update progress bar
        total_loss += loss.item()

    if (batch_idx + 1) % 100 == 0:
        avg_loss = total_loss / (batch_idx + 1)
        progress_bar.set_postfix({'Loss': f'{avg_loss:.4f}'})

```

```

    return total_loss / len(loader)

# Lists to store metrics
train_losses = []

# Training loop with visualization
print(f"Training on {DEVICE}")

for epoch in range(EPOCHS):
    # Train for one epoch
    avg_loss = train_epoch(model, train_loader, optimizer, ddpm_sched,
        NUM_Timesteps)
    train_losses.append(avg_loss)

    # Print progress
    print(f"Epoch {epoch+1}/{EPOCHS} - Average Loss: {avg_loss:.4f}")

    # Generate samples every 10 epochs
    if (epoch + 1) % 10 == 0:
        print(f"\nGenerating samples after epoch {epoch+1}")
        model.eval()
        with torch.no_grad():
            # Generate samples
            samples = ddpm_sample(
                model,
                ddpm_sched,
                img_wh=(28, 28),
                num_ts=NUM_Timesteps,
                seed=epoch
            )

            # Visualize samples
            plt.figure(figsize=(10, 10))
            for i in range(min(16, samples.shape[0])):
                plt.subplot(4, 4, i+1)
                plt.imshow(samples[i, 0].cpu(), cmap='gray')
                plt.axis('off')
            plt.tight_layout()
            plt.show()

# Plot final loss curve
plt.figure(figsize=(10, 5))
plt.plot(train_losses, label='Training Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training Loss Over Time')
plt.legend()

```

```
plt.grid(True)
plt.show()
```

Training on mps

Training: 0%| | 0/1875 [00:00<?, ?it/s]

Epoch 1/20 - Average Loss: 0.1438

Training: 0%| | 0/1875 [00:00<?, ?it/s]

Epoch 2/20 - Average Loss: 0.0742

Training: 0%| | 0/1875 [00:00<?, ?it/s]

Epoch 3/20 - Average Loss: 0.0662

Training: 0%| | 0/1875 [00:00<?, ?it/s]

Epoch 4/20 - Average Loss: 0.0618

Training: 0%| | 0/1875 [00:00<?, ?it/s]

Epoch 5/20 - Average Loss: 0.0589

Training: 0%| | 0/1875 [00:00<?, ?it/s]

Epoch 6/20 - Average Loss: 0.0567

Training: 0%| | 0/1875 [00:00<?, ?it/s]

Epoch 7/20 - Average Loss: 0.0555

Training: 0%| | 0/1875 [00:00<?, ?it/s]

Epoch 8/20 - Average Loss: 0.0542


Training: 0%| | 0/1875 [00:00<?, ?it/s]

Epoch 9/20 - Average Loss: 0.0533

Training: 0%| | 0/1875 [00:00<?, ?it/s]

Epoch 10/20 - Average Loss: 0.0525



Generating samples after epoch 10



```
Training:  0%|          | 0/1875 [00:00<?, ?it/s]
Epoch 11/20 - Average Loss: 0.0517
Training:  0%|          | 0/1875 [00:00<?, ?it/s]
Epoch 12/20 - Average Loss: 0.0512
Training:  0%|          | 0/1875 [00:00<?, ?it/s]
Epoch 13/20 - Average Loss: 0.0507
Training:  0%|          | 0/1875 [00:00<?, ?it/s]
Epoch 14/20 - Average Loss: 0.0503
Training:  0%|          | 0/1875 [00:00<?, ?it/s]
Epoch 15/20 - Average Loss: 0.0500
Training:  0%|          | 0/1875 [00:00<?, ?it/s]
Epoch 16/20 - Average Loss: 0.0494
Training:  0%|          | 0/1875 [00:00<?, ?it/s]
Epoch 17/20 - Average Loss: 0.0494
Training:  0%|          | 0/1875 [00:00<?, ?it/s]
Epoch 18/20 - Average Loss: 0.0493
Training:  0%|          | 0/1875 [00:00<?, ?it/s]
Epoch 19/20 - Average Loss: 0.0488
Training:  0%|          | 0/1875 [00:00<?, ?it/s]
```

Epoch 20/20 - Average Loss: 0.0484

Generating samples after epoch 20



```
[20]: # Generate final test samples
print("\nGenerating final samples")
model.eval()
with torch.no_grad():
    # Generate multiple samples
    samples = []
```

```


for i in range(4): # Generate 4 batches of samples
    sample = ddpm_sample(
        model,
        ddpm_sched,
        img_wh=(28, 28),
        num_ts=NUM_TIMESTEPS,
        seed=i
    )
    samples.append(sample)
samples = torch.cat(samples, dim=0)

# Create a grid of samples
plt.figure(figsize=(15, 15))
for i in range(min(25, samples.shape[0])):
    plt.subplot(5, 5, i+1)
    plt.imshow(samples[i, 0].cpu(), cmap='gray')
    plt.axis('off')
plt.suptitle('Generated Samples')
plt.tight_layout()
plt.show()

# Save the model
torch.save(model.state_dict(), 'time_conditional_unet.pth')
print("Model saved as time_conditional_unet.pth")

```

Generating final samples



Model saved as time_conditional_unet.pth

```

[21]: import random
# Load the saved model
model = TimeConditionalUNet(in_channels=1, num_classes=10, num_hiddens=NUM_HIDDENS).to(DEVICE)
model.load_state_dict(torch.load('time_conditional_unet.pth'))

```

```

# Generate new samples
with torch.no_grad():
    samples = ddpm_sample(
        model,
        ddpm_sched,
        img_wh=(28, 28),
        num_ts=NUM_TIMESTEPS,
        seed=random.randint(0, 1000) # Random seed for reproducibility
    )

# Visualize
plt.figure(figsize=(10, 10))
for i in range(min(16, samples.shape[0])):
    plt.subplot(4, 4, i+1)
    plt.imshow(samples[i, 0].cpu(), cmap='gray')
    plt.axis('off')
plt.tight_layout()
plt.show()

```



3.3.1 Implementing class-conditioned UNet

```

[22]: class ClassConditionalUNet(nn.Module):
    def __init__(
        self,
        in_channels: int,
        num_classes: int,
        num_hiddens: int,
    ):
        super().__init__()
        self.in_channels = in_channels

```



```

self.num_classes = num_classes
self.num_hiddens = num_hiddens

# Time embedding
time_embed_dim = num_hiddens * 4
self.time_embed = nn.Sequential(
    FCBlock(1, time_embed_dim),
    FCBlock(time_embed_dim, time_embed_dim)
)

# Class embedding
self.class_embed = nn.Embedding(num_classes, time_embed_dim)

# Combined embedding projection layers
self.combined_proj_1 = nn.Linear(time_embed_dim * 2, num_hiddens)
self.combined_proj_2 = nn.Linear(time_embed_dim * 2, num_hiddens)
self.combined_proj_3 = nn.Linear(time_embed_dim * 2, num_hiddens * 2)

# Encoder path
self.first_conv = ConvBlock(in_channels, num_hiddens)
self.down1 = DownBlock(num_hiddens, num_hiddens)
self.down2 = DownBlock(num_hiddens, num_hiddens * 2)

# Bottleneck
self.flatten = Flatten()
self.unflatten = Unflatten(num_hiddens * 2)

# Decoder path
self.upsample1 = UpBlock(num_hiddens * 4, num_hiddens)
self.upsample2 = UpBlock(num_hiddens * 2, num_hiddens)

# Final layers
self.final_conv1 = ConvBlock(num_hiddens * 2, num_hiddens)
self.final_conv2 = nn.Conv2d(num_hiddens, in_channels, kernel_size=3,
padding=1)

def forward(
    self,
    x: torch.Tensor,
    c: torch.Tensor,
    t: torch.Tensor,
    mask: torch.Tensor | None = None,
) -> torch.Tensor:
    """
    Args:
        x: (N, C, H, W) input tensor.
        c: (N,) int64 condition tensor.

```

```

        t: (N,) normalized time tensor.
        mask: (N,) mask tensor. If not None, mask out condition when mask_
    == 0.

    Returns:
        (N, C, H, W) output tensor.
    """
    # assert x.shape[-2:] == (28, 28), "Expect input shape to be (28, 28)."

    # Time embedding
    t = t.unsqueeze(-1).float() # (N, 1)
    t_emb = self.time_embed(t) # (N, time_embed_dim)

    # Class embedding
    c_emb = self.class_embed(c) # (N, time_embed_dim)

    # Apply masking if provided
    if mask is not None:
        mask = mask.unsqueeze(-1) # (N, 1)
        c_emb = c_emb * mask # Zero out embeddings for masked indices

    # Combine time and class embeddings
    combined_emb = torch.cat([t_emb, c_emb], dim=1) # (N, time_embed_dim * 2)

    # Project combined embeddings for different layers
    emb1 = self.combined_proj_1(combined_emb).unsqueeze(-1).unsqueeze(-1) # (N, num_hiddens, 1, 1)
    emb2 = self.combined_proj_2(combined_emb).unsqueeze(-1).unsqueeze(-1) # (N, num_hiddens, 1, 1)
    emb3 = self.combined_proj_3(combined_emb).unsqueeze(-1).unsqueeze(-1) # (N, num_hiddens * 2, 1, 1)

    # Encoder
    x1 = self.first_conv(x)
    x1 = x1 + emb1.expand_as(x1)

    x2 = self.down1(x1)
    x2 = x2 + emb2.expand_as(x2)

    x3 = self.down2(x2)
    x3 = x3 + emb3.expand_as(x3)

    # Bottleneck
    x_flat = self.flatten(x3)
    x_bottle = self.unflatten(x_flat)

    # Decoder with skip connections

```

```

x = torch.cat([x_bottle, x3], dim=1)
x = self.upsample1(x)

x = torch.cat([x, x2], dim=1)
x = self.upsample2(x)

x = torch.cat([x, x1], dim=1)
x = self.final_conv1(x)
x = self.final_conv2(x)

return x

```

```

[23]: def ddp_forward(
    unet: ClassConditionalUNet,
    ddpm_schedule: dict,
    x_0: torch.Tensor,
    c: torch.Tensor,
    p_uncond: float,
    num_ts: int,
) -> torch.Tensor:
    """Algorithm 1 of the DDPM paper with classifier-free guidance."""
    unet.train()

    # Sample timestep uniformly
    t = torch.randint(0, num_ts, (x_0.shape[0],), device=x_0.device)

    epsilon = torch.randn_like(x_0)
    alpha_bar = ddpm_schedule['alpha_bars'][t].view(-1, 1, 1, 1)
    x_t = torch.sqrt(alpha_bar) * x_0 + torch.sqrt(1 - alpha_bar) * epsilon

    mask = torch.rand(x_0.shape[0], device=x_0.device) > p_uncond
    epsilon_theta = unet(x_t, c, t / num_ts, mask)
    loss = F.mse_loss(epsilon_theta, epsilon)

    return loss

@torch.inference_mode()
def ddp_sample(
    unet: ClassConditionalUNet,
    ddpm_schedule: dict,
    c: torch.Tensor,
    img_wh: tuple[int, int],
    num_ts: int,
    guidance_scale: float = 5.0,
    seed: int = 0,
) -> torch.Tensor:
    """Algorithm 2 of the DDPM paper with classifier-free guidance."""

```

```

UNET.eval()

torch.manual_seed(seed)

x_t = torch.randn(c.shape[0], unet.in_channels, *img_wh, device=c.device)

for t in reversed(range(num_ts)):
    t_tensor = torch.tensor([t] * c.shape[0], device=c.device)
    t_scaled = t_tensor / num_ts

    eps_cond = unet(x_t, c, t_scaled, mask=torch.ones_like(c, dtype=torch.
→bool))
    eps_uncond = unet(x_t, c, t_scaled, mask=torch.zeros_like(c,
→dtype=torch.bool))

    epsilon_theta = eps_uncond + guidance_scale * (eps_cond - eps_uncond)

    alpha = ddpm_schedule['alphas'][t]
    alpha_bar = ddpm_schedule['alpha_bars'][t]
    beta = ddpm_schedule['betas'][t]

    # Add noise if t > 0
    if t > 0:
        noise = torch.randn_like(x_t)
    else:
        noise = 0

    # Denoising step
    x_t = (1 / torch.sqrt(alpha)) * (
        x_t - (beta / torch.sqrt(1 - alpha_bar)) * epsilon_theta
    ) + torch.sqrt(beta) * noise

return x_t

```

```

[24]: @torch.inference_mode()
def ddpm_sample(
    unet: ClassConditionalUNet,
    ddpm_schedule: dict,
    c: torch.Tensor,
    img_wh: tuple[int, int],
    num_ts: int,
    guidance_scale: float = 5.0,
    seed: int = 0,
) -> torch.Tensor:
    """Algorithm 2 of the DDPM paper with classifier-free guidance."""
    UNET.eval()

```

```

torch.manual_seed(seed)

x_t = torch.randn(c.shape[0], unet.in_channels, *img_wh, device=c.device)

for t in reversed(range(num_ts)):
    t_tensor = torch.tensor([t] * c.shape[0], device=c.device)
    t_scaled = t_tensor / num_ts

    eps_cond = unet(x_t, c, t_scaled, mask=torch.ones_like(c, dtype=torch.
↪bool))
    eps_uncond = unet(x_t, c, t_scaled, mask=torch.zeros_like(c,
↪dtype=torch.bool))

    epsilon_theta = eps_uncond + guidance_scale * (eps_cond - eps_uncond)

    alpha = ddpm_schedule['alphas'][t]
    alpha_bar = ddpm_schedule['alpha_bars'][t]
    beta = ddpm_schedule['betas'][t]

    if t > 0:
        noise = torch.randn_like(x_t)
    else:
        noise = 0

    # Denoising step
    x_t = (1 / torch.sqrt(alpha)) * (
        x_t - (beta / torch.sqrt(1 - alpha_bar)) * epsilon_theta
    ) + torch.sqrt(beta) * noise

return x_t

```

```

[25]: class DDPM(nn.Module):
    def __init__(
        self,
        unet: ClassConditionalUNet,
        betas: tuple[float, float] = (1e-4, 0.02),
        num_ts: int = 300,
        p_uncond: float = 0.1,
    ):
        super().__init__()
        self.unet = unet
        self.betas = betas
        self.num_ts = num_ts
        self.p_uncond = p_uncond
        self.ddpm_schedule = nn.ParameterDict(ddpm_schedule(betas[0], betas[1],
↪num_ts))

```

```

def forward(self, x: torch.Tensor, c: torch.Tensor) -> torch.Tensor:
    """
    Args:
        x: (N, C, H, W) input tensor.
        c: (N,) int64 condition tensor.

    Returns:
        (,) diffusion loss.
    """
    return ddpm_forward(
        self.unet, self.ddpm_schedule, x, c, self.p_uncond, self.num_ts
    )

@torch.inference_mode()
def sample(
    self,
    c: torch.Tensor,
    img_wh: tuple[int, int],
    guidance_scale: float = 5.0,
    seed: int = 0,
):
    return ddpm_sample(
        self.unet, self.ddpm_schedule, c, img_wh, self.num_ts,
        ↪guidance_scale, seed
    )

```

3.4 Training the Class-conditioned UNet

- Plot the loss curve
- Sample results on the test set

```

[26]: # Setup constants and model
EPOCHS = 20
BATCH_SIZE = 128
LEARNING_RATE = 2e-4
NUM_HIDDENS = 64
NUM_TIMESTEPS = 1000
BETA1, BETA2 = 1e-4, 0.02
P_UNCOND = 0.1 # Probability of unconditional training
GUIDANCE_SCALE = 3.0 # For sampling
DEVICE = torch.device("cuda" if torch.cuda.is_available() else "mps" if torch.
    ↪backends.mps.is_available() else "cpu")

# Initialize model and optimizer
model = ClassConditionalUNet(
    in_channels=1,
    num_classes=10, # For MNIST digits 0-9

```

```

        num_hiddens=NUM_HIDDENS
    ).to(DEVICE)

optimizer = optim.Adam(model.parameters(), lr=LEARNING_RATE)

# Create DDPM schedule
ddpm_sched = ddpm_schedule(BETA1, BETA2, NUM_TIMESTEPS)
for k, v in ddpm_sched.items():
    ddpm_sched[k] = v.to(DEVICE)

# Setup data
train_dataset = MNIST(root='data', train=True, transform=ToTensor(),
    ↪download=True)
test_dataset = MNIST(root='data', train=False, transform=ToTensor(),
    ↪download=True)

train_loader = DataLoader(train_dataset, batch_size=BATCH_SIZE, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=BATCH_SIZE, shuffle=False)

```

```

[27]: # Training loop
def train_epoch(model, loader, optimizer, ddpm_sched, num_ts, p_uncond):
    model.train()
    total_loss = 0
    progress_bar = tqdm(loader, desc='Training')

    for batch_idx, (images, labels) in enumerate(progress_bar):
        images = images.to(DEVICE)
        labels = labels.to(DEVICE)

        optimizer.zero_grad()

        # Forward pass
        loss = ddpm_forward(model, ddpm_sched, images, labels, p_uncond, num_ts)

        # Backward pass
        loss.backward()
        optimizer.step()

        # Update progress bar
        total_loss += loss.item()
        avg_loss = total_loss / (batch_idx + 1)
        progress_bar.set_postfix({'Loss': f'{avg_loss:.4f}'})

    return total_loss / len(loader)

```

```

[28]: # Lists to store metrics
train_losses = []

```

```

# Training loop with visualization
print(f"Training on {DEVICE}")

for epoch in range(EPOCHS):
    # Train for one epoch
    avg_loss = train_epoch(model, train_loader, optimizer, ddpm_sched,
        NUM_Timesteps, P_UNCOND)
    train_losses.append(avg_loss)

    # Print progress
    print(f"Epoch {epoch+1}/{EPOCHS} - Average Loss: {avg_loss:.4f}")

    # Generate samples every 10 epochs
    if (epoch + 1) % 10 == 0:
        print(f"\nGenerating samples after epoch {epoch+1}")
        model.eval()
        with torch.no_grad():
            # Generate samples for each digit
            samples_per_digit = 4
            all_samples = []

            for digit in range(10):
                # Create condition tensor for this digit
                c = torch.full((samples_per_digit,), digit, device=DEVICE)

                # Generate samples
                samples = ddpm_sample(
                    model,
                    ddpm_sched,
                    c,
                    img_wh=(28, 28),
                    num_ts=NUM_Timesteps,
                    guidance_scale=GUIDANCE_SCALE,
                    seed=epoch
                )
                all_samples.append(samples)

            # Concatenate all samples
            all_samples = torch.cat(all_samples, dim=0)

            # Visualize samples
            plt.figure(figsize=(15, 6))
            for i in range(40): # Show 4 samples for each digit
                plt.subplot(10, 4, i+1)
                plt.imshow(all_samples[i, 0].cpu(), cmap='gray')
                plt.axis('off')

```



```

        if i % 4 == 0:
            plt.title(f'Digit {i//4}')
            plt.tight_layout()
            plt.show()

# Plot final loss curve
plt.figure(figsize=(10, 5))
plt.plot(train_losses, label='Training Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training Loss Over Time')
plt.legend()
plt.grid(True)
plt.show()

```

Training on mps

```

Training:  0%|          | 0/469 [00:00<?, ?it/s]
Epoch 1/20 - Average Loss: 0.0812
Training:  0%|          | 0/469 [00:00<?, ?it/s]
Epoch 2/20 - Average Loss: 0.0289
Training:  0%|          | 0/469 [00:00<?, ?it/s]
Epoch 3/20 - Average Loss: 0.0249
Training:  0%|          | 0/469 [00:00<?, ?it/s]
Epoch 4/20 - Average Loss: 0.0223
Training:  0%|          | 0/469 [00:00<?, ?it/s]
Epoch 5/20 - Average Loss: 0.0213
Training:  0%|          | 0/469 [00:00<?, ?it/s]
Epoch 6/20 - Average Loss: 0.0204
Training:  0%|          | 0/469 [00:00<?, ?it/s]
Epoch 7/20 - Average Loss: 0.0201
Training:  0%|          | 0/469 [00:00<?, ?it/s]
Epoch 8/20 - Average Loss: 0.0196
Training:  0%|          | 0/469 [00:00<?, ?it/s]
Epoch 9/20 - Average Loss: 0.0189
Training:  0%|          | 0/469 [00:00<?, ?it/s]
Epoch 10/20 - Average Loss: 0.0190

```

Generating samples after epoch 10

Digit 0				
Digit 1				
Digit 2				
Digit 3				
Digit 4				
Digit 5				
Digit 6				
Digit 7				
Digit 8				
Digit 9				

Training: 0%| | 0/469 [00:00<?, ?it/s]

Epoch 11/20 - Average Loss: 0.0183

Training: 0%| | 0/469 [00:00<?, ?it/s]

Epoch 12/20 - Average Loss: 0.0182

Training: 0%| | 0/469 [00:00<?, ?it/s]

Epoch 13/20 - Average Loss: 0.0181

Training: 0%| | 0/469 [00:00<?, ?it/s]

Epoch 14/20 - Average Loss: 0.0175

Training: 0%| | 0/469 [00:00<?, ?it/s]

Epoch 15/20 - Average Loss: 0.0177

Training: 0%| | 0/469 [00:00<?, ?it/s]

Epoch 16/20 - Average Loss: 0.0176

Training: 0%| | 0/469 [00:00<?, ?it/s]

Epoch 17/20 - Average Loss: 0.0171

Training: 0%| | 0/469 [00:00<?, ?it/s]

Epoch 18/20 - Average Loss: 0.0174

Training: 0%| | 0/469 [00:00<?, ?it/s]


Epoch 19/20 - Average Loss: 0.0173

Training: 0%| | 0/469 [00:00<?, ?it/s]

Epoch 20/20 - Average Loss: 0.0166

Generating samples after epoch 20

Digit 0	0	0	0	0
Digit 1	1	1	1	1
Digit 2	2	2	2	2
Digit 3	3	3	3	3
Digit 4	4	4	4	4
Digit 5	5	5	5	5
Digit 6	6	6	6	6
Digit 7	7	7	7	7
Digit 8	8	8	8	8
Digit 9	9	9	9	9



```
[29]: # Save the model
torch.save(model.state_dict(), 'class_conditional_unet.pth')
print("Model saved as class_conditional_unet.pth")
```

Model saved as class_conditional_unet.pth

```
[30]: # Load the saved class-conditional UNet model
model = ClassConditionalUNet(
    in_channels=1,
    num_classes=10,
    num_hiddens=NUM_HIDDENS
).to(DEVICE)
model.load_state_dict(torch.load('class_conditional_unet.pth',
    ↪map_location=DEVICE))
model.eval()
print("Model loaded from class_conditional_unet.pth")
```

Model loaded from class_conditional_unet.pth

```
[31]: # Generate final test samples
print("\nGenerating final samples")
model.eval()
with torch.no_grad():
    # Generate multiple samples for each digit
    samples_per_digit = 5
    all_samples = []

    for digit in range(10):
        # Create condition tensor for this digit
        c = torch.full((samples_per_digit,), digit, device=DEVICE)

        # Generate samples with different guidance scales
        for scale in [1.0, 3.0, 5.0]:
            samples = ddpm_sample(
                model,
                ddpm_sched,
                c,
                img_wh=(28, 28),
                num_ts=NUM_TIMESTEPS,
                guidance_scale=scale,
                seed=digit
            )
            all_samples.append(samples)

    # Concatenate all samples
    all_samples = torch.cat(all_samples, dim=0)


    # Create a grid of samples
```

```

plt.figure(figsize=(15, 10))
for i in range(len(all_samples)):
    plt.subplot(10, 15, i+1)
    plt.imshow(all_samples[i, 0].cpu(), cmap='gray')
    plt.axis('off')
    if i % 15 == 0:
        plt.title(f'Digit {i//15}', y=1.0)
plt.suptitle('Generated Samples (Columns: Different Guidance Scales)', y=0.
↪95)
plt.tight_layout()
plt.show()

```

Generating final samples



```

[32]: # Function to generate samples for a specific digit
def generate_digit_samples(model, digit, num_samples=4, guidance_scale=3.0):
    model.eval()
    with torch.no_grad():
        c = torch.full((num_samples,), digit, device=DEVICE)
        samples = ddpm_sample(
            model,
            ddpm_sched,
            c,

```

```

        img_wh=(28, 28),
        num_ts=NUM_TIMESTEPS,
        guidance_scale=guidance_scale
    )


    plt.figure(figsize=(10, 3))
    for i in range(num_samples):
        plt.subplot(1, num_samples, i+1)
        plt.imshow(samples[i, 0].cpu(), cmap='gray')
        plt.axis('off')
    plt.suptitle(f'Generated Samples for Digit {digit}')
    plt.show()

# Example usage:
print("\nGenerating samples for specific digits:")
for digit in [0, 1, 4, 7, 9]:
    generate_digit_samples(model, digit)


```

Generating samples for specific digits:


Generated Samples for Digit 0




Generated Samples for Digit 1



Generated Samples for Digit 4



Generated Samples for Digit 7



Generated Samples for Digit 9

